

Implementing RNN, LSTM, Sequence and Transformer

1. Implementing RNN


You should make sure you have the following packages installed

```
$ pip install torchtext==0.8.1
$ pip install torch==1.7.1
$ pip install spacy==2.3.5
$ pip install tqdm
$ pip install numpy
```

Additionally, you will need the Spacy tokenizers in English and German language, which can be downloaded as such:

```
$python -m spacy download en
$python -m spacy download de
```

You will be using PyTorch Linear layers and activations to implement a vanilla RNN unit. Please refer to the following structure and complete the code in `RNN.py`:



2. Implement LSTM

You will be using PyTorch `nn.Parameter` and activations to implement an LSTM unit. You can simply translate the following equations using `nn.Parameter` and PyTorch activation functions to build an LSTM from scratch:

$$\begin{aligned}i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{t-1} + b_{hi}) \\f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{t-1} + b_{hf}) \\g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{t-1} + b_{hg}) \\o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{t-1} + b_{ho}) \\c_t &= f_t \odot c_{t-1} + i_t \odot g_t \\h_t &= o_t \odot \tanh(c_t)\end{aligned}$$

Here's a great visualization of the above equation to help you understand LSTM unit.
<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>


3. Seq2Seq Implementation

`seq2seq.py` you will see the files needed to complete this section. In these files you will complete the initialization and forward pass in `__init__` and forward function. `Encoder.py` `Decoder.py` `Seq2Seq.py`.

Train `seq2seq` on the dataset with the default hyperparameters.

4. Transformers Implementation

We will be implementing a one-layer Transformer encoder which, similar to an RNN, can encode a sequence of inputs and produce a final output of possibility of tokens in target language. The architecture can be seen below. You will see the file `Transformer.py`. You will implement the functions in the `TransformerTranslator` class.



You can refer to the original paper in the link below:

<https://arxiv.org/pdf/1706.03762.pdf> for more details

Additional help for Transformer Implementation

We will format our input embeddings similarly to how they are constructed in [BERT (source of figure)](<https://arxiv.org/pdf/1810.04805.pdf>). Unlike a RNN, a Transformer does not include any positional information about the order in which the words in the sentence occur. Because of this, we need to append a positional encoding token at each position. (We will ignore the segment embeddings and [SEP] token here, since we are only encoding one sentence at a time). We have already appended the [CLS] token for you in the previous step.

Input	[CLS]	my	dog	is	cute	[SEP]	he	likes	play	#ing	[SEP]
Token Embeddings	$E_{[CLS]}$	E_{my}	E_{dog}	E_{is}	E_{cute}	$E_{[SEP]}$	E_{he}	E_{likes}	E_{play}	E_{ing}	$E_{[SEP]}$
	+	+	+	+	+	+	+	+	+	+	+
Segment Embeddings	E_A	E_A	E_A	E_A	E_A	E_A	E_B	E_B	E_B	E_B	E_B
	+	+	+	+	+	+	+	+	+	+	+
Position Embeddings	E_0	E_1	E_2	E_3	E_4	E_5	E_6	E_7	E_8	E_9	E_{10}

Your first task is to implement the embedding lookup, including the addition of positional encodings. Complete the code section for Deliverable 1, which will include part of `__init__` and `embed`.

Attention can be computed in matrix-form using the following formula:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

We want to have multiple self-attention operations, computed in parallel. Each of these is called a *head*. We concatenate the heads and multiply them with the matrix *attention_head_projection* to produce the output of this layer.

After every multi-head self-attention and feedforward layer, there is a residual connection + layer normalization. Make sure to implement this, using the following formula:

$$\text{LayerNorm}(x + \text{Sublayer}(x))$$

Implement the function `multi_head_attention` for **Deliverable 2**. We have already initialized all of the layers you will need in the constructor.

Complete code for **Deliverable 3** in `feedforward_layer`: the element-wise feed-forward layer consisting of two linear transformers with a ReLU layer in between.

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

Complete code for **Deliverable 4** in `final_layer`., to produce probability scores for all tokens in target language.

Put it all together by completing the method `forward`, where you combine all of the methods you have developed in the right order to perform a full forward pass.

Train the transformer architecture on the dataset with the default hyperparameters – you should get a perplexity better than that for seq2seq. Then perform hyperparameter tuning and include the improved results.

Include the accuracy of the Seq2Seq model and Transformer architecture before and after hyperparameter tuning.